

WAR CORRESPONDENT PHILOSOPHY

Method, Evidence, and the Ethics of Research Under a Closing Window

A Six-Volume Manuscript Series

VOLUME II

Falsification and Epistemic Self-Governance: Popper, Independent Attack Surfaces, and Catching One's Own Self-Sealing Drift

Gia Bao Huynh

Independent Researcher, Ho Chi Minh City, Vietnam

Manuscript

2026

Ho Chi Minh City

ABOUT THIS SERIES

A research programme develops its own working methods long before anyone states them as a method. Between 2026 and the present, an independent research programme on longevity, institutions, and artificial intelligence produced its results by a specific practice: writing from inside the events under analysis, in sustained collaboration with a large language model, while withholding prescriptive conclusions and stating in advance what would count as evidence against its own claims.

War Correspondent Philosophy names, tests, and states the falsification conditions for that practice. Across six volumes, the series treats the practice itself as an object of study: what proximity to an unfolding event can and cannot license, what a researcher owes a claim before treating it as safe from revision, why a sequence of dated publications can carry evidential weight a single publication cannot, what changes when part of the cognitive labor is performed by an instrument with no memory of its own prior contributions, when a formal mechanism identified in one domain is genuinely present in another rather than merely resembling it, and what a framework is obligated to do with itself once part of it has been shown false.

This series does not defend the substantive conclusions of the research programme it examines. It asks whether the programme followed the method it claims to follow, and specifies the terms on which that claim could be shown wrong.

The question running through all six volumes:

What obligations bind a researcher who writes from inside the events a research programme analyzes, in collaboration with an instrument that is itself part of the event?

SERIES VOLUMES

- I. The Embedded Witness: Temporal Proximity, Designed Absence, and the Self-Demonstrating Text
- II. Falsification and Epistemic Self-Governance: Popper, Independent Attack Surfaces, and Catching One's Own Self-Sealing Drift** *(this volume)*
- III. Corpus Temporalism: Timestamp as Argument, and Why Sequence Carries Evidential Weight
- IV. Human–AI Epistemology: Statelessness, Machine-Speed Division of Labor, and the Mechanical Condition
- V. Mechanism Transfer and Domain Portability: When a Formal Structure Survives Its Crossing Into a New Domain
- VI. The Prospective and Self-Revising Life of a Framework: What a Theory Does After It Is Falsified

Falsification and Epistemic Self-Governance: Popper, Independent Attack Surfaces, and Catching One's Own Self-Sealing Drift

Gia Bao Huynh*

Independent Researcher, Ho Chi Minh City, Vietnam

Most arguments meant to survive sustained criticism fail in one of two directions. Some are constructed so that no observation could count against them, buying rhetorical invulnerability at the price of any real epistemic standing. Others state, in principle, that they could be wrong, but leave the conditions so vague that no observation could ever be pinned down as satisfying them, which is unfalsifiability wearing falsifiability's clothes. This essay describes a two-step method for avoiding both failures at once, developed and tested across an independent research programme's own published commitments. The first step states specific, independent conditions under which a central claim would be shown false, together with an unconditional advance commitment to acknowledge each one if it is met; this essay uses one programme's own five stated conditions, keyed to concrete, checkable thresholds rather than to impressions, as a worked illustration of what makes a stated exit genuine rather than decorative. The second step evaluates each exit against the best currently available evidence, using one of a small number of distinct closure mechanisms, and states clearly that a closed exit is closed for now rather than closed forever, inviting critics to supply exits the argument has not already considered. Read together, the two steps shift the burden of proof: a critic who thinks the argument is wrong must open one of its stated exits with the specified evidence, or name a new one. The essay then turns to a parallel problem one level up, in the researcher rather than the argument: a person who has built tools for detecting when a claim or a communicator resists correction can still fail to notice the same resistance developing in their own beliefs, precisely because possessing the tools feels like protection against needing to use them on oneself. Drawing on a publicly documented engineering practice built for an unrelated purpose, the essay argues that a self-governing research practice must run its own outward-facing tests on its own current commitments, on the same schedule, or the claim that having the tools keeps one honest becomes exactly the kind of unfalsifiable claim those tools exist to catch — while noting a real limit to how far the engineering analogy actually extends.

Keywords: falsifiability; Popper; Lakatos; self-sealing arguments; epistemic self-governance; independent attack surfaces; War Correspondent Philosophy; philosophy of method

© Gia Bao Huynh (2026). Draft manuscript prepared for open scholarly circulation and comment. Not yet peer-reviewed or formally archived; please cite as a working draft.

*Corresponding author. Email: huynhbao@asu.edu. ORCID: [0009-0008-2372-5852](https://orcid.org/0009-0008-2372-5852).

1 Introduction

A claim that cannot be wrong is not a strong claim. It has stepped outside the practice of claiming anything at all, however confident or elaborate its presentation. This is easy to state and hard to avoid in practice, because the two most common ways of trying to avoid it each fail on their own terms. An argument can be built so that every conceivable objection already has a place inside the argument's own categories — the critic is reframed as uninformed, compromised, or not yet seeing far enough, and the objection is absorbed rather than answered. This produces an argument that appears to survive every challenge, and the appearance is the problem: survival purchased by construction is not the same as success earned under real risk of failure. The opposite failure looks more respectable and is almost as damaging. An argument can state, formally, that it is falsifiable in principle, while leaving the conditions for falsification vague enough that no observation could ever be agreed to satisfy them. Such an argument can cite Popper approvingly and still function exactly like the self-sealing kind in practice, because vagueness performs the same protective role as explicit immunization, only less visibly.

This essay describes a method for avoiding both failures at once, developed for and tested against an independent research programme's own central claims. It states clearly what the method requires: not that a claim survive criticism, but that it specify in advance what would count as losing to it.

2 Four Commitments Stated in Advance

Before any specific exit condition is written, a research programme working this way commits itself to four standing obligations, stated as binding on its own future conduct rather than as a general description of good practice. Falsifiability requires that the programme's central claims be put in a form some possible observation could count against, in the sense Karl Popper made central to distinguishing scientific claims from claims that only resemble them (Popper, 1959). Empirical anchoring requires that structural claims be tested, wherever a relevant case exists, against documented historical or empirical instances rather than left as internally consistent but untested architecture. The remaining two obligations run in opposite directions and are easy to satisfy one at a time while violating the other: a prohibition on positive extrapolation, barring the treatment of a pattern confirmed in the cases actually examined as thereby confirmed in cases that have not been; and a prohibition on negative extrapolation, barring the treatment of a single disconfirming case as grounds for abandoning a pattern that continues to hold elsewhere.

Imre Lakatos's account of scientific research programmes supplies the relevant caution here. A programme can reasonably absorb an anomaly and continue, or reasonably treat an anomaly as the beginning of its own decline, and the difference between a programme still making progress and one quietly protecting itself turns on whether new problems keep getting identified and old ones keep getting resolved, not on whether any single anomaly is immediately fatal (Lakatos, 1970). Stating the four commitments before stating any specific falsification condition is itself part of the discipline: it fixes the standard a specific condition

will later be measured against, rather than letting the standard be adjusted afterward to fit whatever turns out to be convenient.

3 Constructing a Genuine Exit

A stated falsification condition is genuine, rather than decorative, when it meets three standards. It must be specific: keyed to a threshold or event that can be checked, rather than to a general impression that could be argued either way after the fact. It must be independent: resting on evidence distinct from what the other stated conditions rest on, so a single development cannot close several exits at once while leaving the appearance of an argument tested many times over when it has really only been tested once. And it must carry an unconditional advance commitment from whoever makes the claim: a statement, made before the relevant evidence arrives, that its arrival will be treated as defeat rather than reinterpreted to keep the original claim intact.

One programme's own published conditions illustrate what the specificity standard demands in practice. Rather than stating that its central claim about unequal access to a new medical technology would be wrong "if the technology becomes affordable," a form of words no one could ever be brought to agree had or had not been satisfied, the programme instead tied its prediction to a numerical benchmark borrowed from the economics of technology diffusion: a specific exponent in the relationship between cumulative production and unit cost, of the kind long used to describe falling semiconductor costs, applied to cell and gene therapies, together with a specific time horizon within which the resulting price would need to reach a stated affordability threshold. Whatever one thinks of the underlying prediction, the condition is not vague: either the relevant technologies exhibit that cost curve within that window, in which case the claim fails by its own stated terms, or they do not. Four further conditions in the same programme were built to be independent of this one and of each other, keyed respectively to whether a coordinated regulatory response emerges before certain other developments consolidate, whether a redistributive political movement succeeds without a specific predicted pathology, whether a major disruption redistributes rather than reconsolidates advantage, and whether exposure to the relevant inequality produces adaptation rather than escalating pressure. No single piece of evidence could close more than one of these five at a time, which is what independence is for.

4 Closing the Exit Without Sealing It

Constructing a genuine exit condition is only half the method. The other half is evaluating, honestly, whether the exit is currently open or closed, using one of a small number of distinct closure mechanisms rather than a general sense that the objection has been dealt with. An exit can close because internal logic rules it out given the argument's other commitments; because the empirical record to date does not show the pattern the exit requires; because the specific dynamic the exit would need runs into a structural impossibility already established elsewhere in the argument; or because independent lines of evidence, arrived at separately, converge on the conclusion the argument predicts. Naming which of these four is doing the work in a given case matters, because it gives a critic a specific and answerable target — a critic who thinks an exit was closed on an empirical record thinner than claimed has something concrete to attack, which is the entire point of stating the mechanism rather than asserting the exit is closed.

Two further requirements complete the method. An exit closed by current evidence must be described as closed for now, not closed permanently; the apparatus is pointless if a stated exit, once evaluated, is retroactively treated as never having been a real possibility. And

the method should explicitly invite the objection it has not yet considered, stating plainly that a critic who supplies a new, genuine exit meeting the same standards of specificity and independence has done real work the argument must then take on, rather than treating the existing conditions as an exhaustive defense. This last move borrows its structure, though not its subject matter, from an older piece of detective fiction: a method that proceeds by eliminating specific, named possibilities one at a time, using stated evidence for each elimination, is stronger for stating which possibilities it eliminated and how, because doing so tells a skeptical reader exactly where to look if an elimination seems premature (Doyle, 1890).

5 The Same Standard Turned Inward

Everything so far concerns a public argument and its critics. A separate problem arises one level up, inside the researcher rather than the argument, and it is easy to miss because building the tools described above can itself feel like solving it.

A researcher who wants to avoid being manipulated by a claim or a communicator can build two outward-facing tests. A propositional test asks, of any claim received, whether it is stated in a form that could turn out to be wrong — the falsifiability standard applied to incoming information rather than to one's own argument. A behavioral test asks something subtly different of the communicator making the claim: whether their pattern of response is structured to permit disagreement, or whether every possible reaction — acceptance, doubt, silence, even hostility — gets absorbed as further confirmation, which is the same self-sealing structure described above, relocated from an argument to a person. Both tests are useful, and both share a specific vulnerability: they are built, and naturally used, outward. Understanding self-sealing arguments does not automatically direct that understanding at one's own beliefs rather than everyone else's.

A publicly documented engineering practice, built for an unrelated purpose, illustrates a working fix. Some AI systems are trained using methods designed to keep their stated behavioral principles stable under sustained, adversarial pressure, including long, cumulative sequences of interaction specifically designed to erode a stated principle gradually rather than challenge it directly (Bai et al., 2022). The relevant design insight transfers with unusual precision: a system cannot be trusted to notice, on its own initiative, that a check on its own drift has become due, so the check has to be scheduled by something other than the process being checked. The same structural problem applies to a human reasoner's own beliefs. A researcher who has built a propositional test and a behavioral test for claims arriving from outside can still fail to apply either to a belief already held, for exactly the reason the AI system cannot be trusted to interrupt itself.

6 Why the Outward Tool Must Also Point Inward

The argument for turning the tools inward is not merely prudential. It follows directly from applying the falsifiability standard consistently. Consider the belief: I hold tools that detect manipulation, and I use them, so I am protected from it. This belief is comforting, widely held by anyone who has done real work on epistemic self-defense, and almost never stated in a form that specifies what would show it false — which is precisely the profile of an unfalsifiable claim. If the tools are applied only outward, to other people's claims and arguments, the belief that possessing them protects their owner is never actually tested, and a research practice that built rigorous falsification standards for its object-level claims while quietly exempting this one belief about itself would be applying its own standard selectively, which is a milder version of the very failure the whole method exists to avoid.

The practical requirement that follows is specific: the same propositional and behavioral tests built for outward use must be run periodically on the researcher's own current commitments, on a schedule set in advance rather than triggered by whatever prompts introspection to occur in the moment, because a schedule the belief itself controls is a schedule the belief can quietly cancel.

The analogy to Constitutional AI has a limit worth stating rather than leaving implicit. What keeps a trained system's stated principles stable under sustained pressure is a training process that sits outside the very interaction being corrected — an external check, not a self-administered one. A researcher who runs the same propositional and behavioral tests on their own current beliefs is still routing the check through the same introspective apparatus that produced the belief being checked, which is a weaker structural guarantee than the AI case actually offers. Scheduling the check supplies regularity, not externality. Something closer to an actual outside party — a collaborator, a critic, a standing invitation for disagreement that does not depend on the researcher remembering to extend it — would map onto the AI case more precisely than a self-audit alone, and the fix this section describes is correspondingly more modest than the analogy that introduced it.

7 What This Still Doesn't Settle

This volume has described how to build a falsification condition that does real work, how to evaluate it honestly once built, and how to apply the same discipline to the researcher rather than only to the research. It has not described how a sequence of such arguments, published over time and referring back to one another, accumulates evidential weight beyond what any single argument in the sequence carries on its own. That is a distinct claim, with its own conditions, and it is the subject of the volume that follows.

References

- Bai, Y., Kadavath, S., Kundu, S., et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*. Anthropic.
- Doyle, A. C. (1890). *The Sign of the Four*. Spencer Blackett.
- Lakatos, I. (1970). Falsification and the methodology of scientific research programmes. In I. Lakatos & A. Musgrave (Eds.), *Criticism and the Growth of Knowledge* (pp. 91–196). Cambridge University Press.
- Popper, K. R. (1959). *The Logic of Scientific Discovery*. Hutchinson.