

# **WAR CORRESPONDENT PHILOSOPHY**

*Method, Evidence, and the Ethics of Research Under a Closing Window*

*A Six-Volume Manuscript Series*

## **VOLUME IV**

### **Human–AI Epistemology: Statelessness, Machine-Speed Division of Labor, and the Mechanical Condition**

Gia Bao Huynh

*Independent Researcher, Ho Chi Minh City, Vietnam*

Manuscript

2026

Ho Chi Minh City

## ABOUT THIS SERIES

A research programme develops its own working methods long before anyone states them as a method. Between 2026 and the present, an independent research programme on longevity, institutions, and artificial intelligence produced its results by a specific practice: writing from inside the events under analysis, in sustained collaboration with a large language model, while withholding prescriptive conclusions and stating in advance what would count as evidence against its own claims.

*War Correspondent Philosophy* names, tests, and states the falsification conditions for that practice. Across six volumes, the series treats the practice itself as an object of study: what proximity to an unfolding event can and cannot license, what a researcher owes a claim before treating it as safe from revision, why a sequence of dated publications can carry evidential weight a single publication cannot, what changes when part of the cognitive labor is performed by an instrument with no memory of its own prior contributions, when a formal mechanism identified in one domain is genuinely present in another rather than merely resembling it, and what a framework is obligated to do with itself once part of it has been shown false.

This series does not defend the substantive conclusions of the research programme it examines. It asks whether the programme followed the method it claims to follow, and specifies the terms on which that claim could be shown wrong.

The question running through all six volumes:

*What obligations bind a researcher who writes from inside the events a research programme analyzes, in collaboration with an instrument that is itself part of the event?*

---

## SERIES VOLUMES

- I. The Embedded Witness: Temporal Proximity, Designed Absence, and the Self-Demonstrating Text
- II. Falsification and Epistemic Self-Governance: Popper, Independent Attack Surfaces, and Catching One's Own Self-Sealing Drift
- III. Corpus Temporalism: Timestamp as Argument, and Why Sequence Carries Evidential Weight
- IV. Human–AI Epistemology: Statelessness, Machine-Speed Division of Labor, and the Mechanical Condition** *(this volume)*
- V. Mechanism Transfer and Domain Portability: When a Formal Structure Survives Its Crossing Into a New Domain
- VI. The Prospective and Self-Revising Life of a Framework: What a Theory Does After It Is Falsified

## **Human–AI Epistemology: Statelessness, Machine-Speed Division of Labor, and the Mechanical Condition**

Gia Bao Huynh\*

*Independent Researcher, Ho Chi Minh City, Vietnam*

A researcher who works with a large language model across an extended inquiry faces a description problem existing vocabularies handle badly. Calling the system a tool undersells outputs that visibly redirect the inquiry's trajectory; calling the relationship collaboration imports assumptions — that the system has opinions, that its persistence with a position reflects commitment, that its agreement counts as corroboration — that are false of current systems and compromise the inquiry when treated as true. This essay develops a third description, built from three specific properties of the participant rather than from a general stance about artificial intelligence. First, the system's characteristic lack of persistence across sessions, usually treated as a nuisance, converts into a specific methodological asset once understood structurally: a participant that begins each session without the accumulated weight of a research programme's prior framings can raise objections a persistent collaborator, gradually shaped by the programme's own history, would stop raising — a benefit this essay argues is real but conditional, and increasingly qualified by memory tooling that partially restores the very continuity the asset depends on the absence of. Second, extending an existing account of the division of cognitive labor in a research community and an existing account of the conditions under which testimony warrants trust shows the resulting division to be genuine but sharply asymmetric: broad, fast, stochastic coverage on one side, and sole responsibility for evaluation, integration, and accountability on the other, because the system's outputs do not meet the conditions that would let them function as trustworthy testimony if accepted uncritically. Third, the essay specifies a mechanical condition with three components — transparency about how outputs are produced, documentation of the interaction that produced them, and calibration of the epistemic weight they are given — and argues that a researcher who honestly treats the system as a characterizable mechanical process integrates its contributions more legitimately than one who treats it as something closer to a human partner.

**Keywords:** human–AI epistemology; statelessness; division of cognitive labor; epistemic dependence; testimony; mechanical condition; War Correspondent Philosophy

---

Huynh, Gia Bao (2026) Human–AI Epistemology: Statelessness, Machine-Speed Division of Labor, and the Mechanical Condition. Volume IV of *War Correspondent Philosophy: Method, Evidence, and the Ethics of Research Under a Closing Window*.

© Gia Bao Huynh (2026). Draft manuscript prepared for open scholarly circulation and comment. Not yet peer-reviewed or formally archived; please cite as a working draft.

---

\*Corresponding author. Email: [huynhbao@asu.edu](mailto:huynhbao@asu.edu). ORCID: [0009-0008-2372-5852](https://orcid.org/0009-0008-2372-5852).

## 1 Introduction

A researcher who asks a large language model to generate competing hypotheses, reconstruct an unfamiliar argument, or identify a counterexample is doing something existing vocabularies describe badly. Calling the exchange the use of a tool captures the researcher’s continuing authority over the inquiry, but strains against outputs that were not fixed in advance and that visibly redirect where the inquiry goes next. Calling the exchange collaboration captures that redirection, but imports, usually without anyone deciding to import it, a cluster of assumptions about the system’s epistemic properties — that it has something like opinions, that its persistence with a position reflects something like commitment, that its agreement functions as something like independent corroboration — that are false of current systems and compromise an inquiry that treats them as true.

This volume develops a third description, built from three specific properties of the participant involved rather than from a general stance on artificial intelligence. None of the three requires settling a difficult question about whether the system really understands anything. All three are visible in how it behaves across a sustained research process, and each carries a specific methodological consequence for the person working with it.

## 2 A Participant With No Memory

The most immediately obvious property is also the one most often treated as a nuisance rather than as data: a large language model does not retain, in any usable sense, its prior contributions to a research programme once a session ends. It begins the next conversation without access to the accumulated reasoning, the rejected framings, the specific objections already raised and answered, unless a human explicitly reintroduces them. Treated only as a limitation, this statelessness is an inconvenience to be engineered around — better notes, re-supplied context, continuity kept in the human’s own files rather than the system’s.

Treated as a structural property of the participant, statelessness does something more specific. A collaborator who accumulates a programme’s history gradually absorbs that programme’s own assumptions and its sense of which objections have already been dealt with, and over enough sessions becomes less likely to raise a genuinely novel objection precisely because it has been shaped by the programme’s own prior responses to objections. A system that begins each session fresh cannot be shaped this way. Provided the researcher does not uncritically re-supply the programme’s prior framing as settled background before asking a question, a stateless system can function as an adversarial generator structurally resistant to the inquiry’s own confirmation tendencies — not because it is more rigorous than a persistent collaborator would be, but because it has not yet been trained, by the specific history of this specific inquiry, to stop raising the objection the inquiry has grown comfortable dismissing (Huynh, 2026a).

This condition is easier to state than to guarantee in practice, and it is worth being honest about exactly where it can fail. Tools that let an AI system carry a research programme’s established vocabulary, prior corrections, and stated preferences forward across sessions — increasingly available, and in active use in research practice generally, including in the production of this series — are close to a direct instance of the re-supplied settled background this section’s own condition warns against. Where such tools are in use, the adversarial value described above is not simply present or absent; it is a matter of degree, tracking how much of the prior framing the tooling actually restores and how critically the researcher treats what it

restores, rather than a property guaranteed by statelessness at the level of the underlying model alone. This converts what looks like the system’s central limitation into a usable asset only conditionally, and the condition is one a researcher using memory-enabled tools has to actively maintain rather than simply inherit.

### **3 Dividing Labor at Machine Speed**

Philip Kitcher’s account of the division of cognitive labor describes how a research community benefits from distributing investigative effort across a problem’s surface rather than converging all at once on whichever approach looks most promising, so the space of live possibilities gets covered even though no individual researcher covers all of it (Kitcher, 1990). Human–AI research collaboration instantiates a version of this division with an unusual asymmetry built in. The AI participant can generate candidate hypotheses and framings across a wide area of a problem’s surface at a speed no individual researcher matches in the same span of time. But this coverage is stochastic rather than reliable, uneven in depth, and unsupported by anything resembling the standing commitment or capacity to distinguish knowledge from confabulation that would let its outputs function as testimony in John Hardwig’s specific sense (Hardwig, 1985). Hardwig’s conditions for warranted trust in testimony require that the party being trusted actually have relevant evidence, that they communicate it reliably, and that the party doing the trusting have good reason to believe both. AI outputs, accepted uncritically, do not reliably satisfy any of the three, a connection developed at length elsewhere (Huynh, 2026b).

The resulting division of labor is genuine, but it assigns the two participants structurally different roles rather than splitting one role in half. The AI occupies the generative and exploratory position: surveying the surface, producing candidate structures, testing the consequences of a proposed distinction quickly enough to try many before committing to one. The human occupies the reflective and evaluative position: selecting which candidates to pursue, assessing their quality against standards the outputs cannot apply to themselves, integrating whatever survives into a developing argument, and bearing sole responsibility for the result. Neither position alone is sufficient for the inquiry, and the division is not a concession to the system’s limitations so much as the structure that makes the collaboration answerable rather than merely fast.

### **4 The Trap on Either Side**

Two failure modes threaten a researcher who does not hold this division firmly in view. Epistemic inflation treats AI outputs as carrying the credibility of testimony from an entity with genuine beliefs, so a conclusion the system produces repeatedly, or with apparent confidence, gets mistaken for a conclusion independently corroborated or well warranted, when repetition and confidence are both features of generated text rather than reliable signals about the underlying claim. Accountability evasion treats the exchange as collaboration in the full human sense, letting the question of responsibility blur, so a researcher who reached a conclusion through extended interaction with the system comes to feel responsibility is somehow shared, when the system bears none of it and the human remains, in every case, the sole party who can be held to account. Avoiding the first failure by falling back on treating the AI as a plain instrument compromises the description in the opposite direction, understating how far its outputs can genuinely redirect an inquiry’s trajectory. The task is not to choose the less inaccurate vocabulary. It is to replace both with a description built from what the participant actually is.

## 5 The Mechanical Condition

That description has three components. Transparency requires understanding, at a level adequate to using the system responsibly, that its outputs are generated by pattern-matching on training data as shaped by the specific prompt, not by a process that tracks evidence or maintains a standing commitment to what it produces. Documentation requires keeping a record of what was asked, what was generated, what was accepted, what was rejected, and what prompted further investigation, because an AI’s contribution is characterizable only by reference to the interaction that produced it, and that record disappears behind the surface of the resulting text unless it is kept deliberately. Calibration requires receiving the system’s contributions at the epistemic level they actually warrant — as hypotheses and candidate framings to be tested, not as conclusions requiring only confirmation.

The counterintuitive result is that satisfying all three makes the collaboration more genuine, not less. A researcher who treats the system as a characterizable mechanical process, with known limitations and a documented interaction structure, integrates its contributions accurately: they are evaluated as what they are, materials for inquiry, by standards appropriate to their source. A researcher who instead attributes to the system something closer to belief or independent judgment builds the inquiry partly on properties it does not have, and the inquiry is not rescued by some of its conclusions turning out to be correct anyway. De-anthropomorphizing the participant is the condition for using it well, not a concession that makes the collaboration less real.

## 6 Why the Record Itself Is Part of the Method

The documentation component is not merely good research hygiene. If the AI’s contributions are genuinely constitutive of an inquiry’s development — if the inquiry would have gone differently without them — the inquiry’s actual development is not fully visible in its published outputs alone, since a finished text records conclusions rather than the sequence of framings tried and discarded on the way to them. A kept record of that sequence is not a transparency courtesy extended to the reader. It is the only available evidence that a given contribution was genuinely constitutive rather than convenient to claim afterward, which connects this volume directly to the corpus-temporalism argument of the previous one: a documented order of interaction is what lets a later reader distinguish an inquiry actually shaped by its stated method from one that adopted the vocabulary of the method after the fact.

## 7 What This Doesn’t Settle

This volume has described three properties of a specific kind of research participant — memorylessness treated as a conditional resource, a genuinely asymmetric division of labor, and a mechanical condition that makes honest use of the participant possible — without addressing a separate question the next volume takes up: whether a formal mechanism identified through work of this kind, in one domain, is ever genuinely present in a different domain, or only resembles it closely enough to tempt the comparison.

## References

- Hardwig, J. (1985). Epistemic dependence. *Journal of Philosophy*, 82(7), 335–349.
- Huynh, G. B. (2026a). *Statelessness: The Engineering, Legal, and Psychological Reasons a Model Forgets Between Sessions, What Forgetting Costs a User and a Decades-Long*

*Research Programme, and the One Epistemic Problem Only Forgetting Can Solve.* Zenodo. <https://doi.org/10.5281/zenodo.21898833>

Huynh, G. B. (2026b). *Epistemic Dependence at Machine Speed: Extending Kitcher's Division of Cognitive Labor and Hardwig's Conditions for Warranted Trust to Human–AI Research Collaboration.* Zenodo. <https://doi.org/10.5281/zenodo.21898900>

Kitcher, P. (1990). The division of cognitive labor. *Journal of Philosophy*, 87(1), 5–22.